


UNIT I: Introduction

Contents:

Definition and scope of Generative AI Overview of generative models and their applications Importance of Generative AI in various domains. AI vs ML vs Generative AI, Uses Self Supervised Learning, Life cycle of GenAI, Exploring Predictive Machine Learning vs Generative AI, Explainable AI

Model Performance vs. Data



- *Small model performance plateaus after a certain volume of data.*
- *Large model performance continues to improve with the volume of data.*

Definition and Scope of Generative AI

Generative AI refers to a subset of artificial intelligence focused on creating new content that mimics human creativity and intelligence. Unlike traditional AI models that classify or predict based on existing data, generative AI models produce new data instances. This capability is pivotal in fields such as creative industries, healthcare, and finance, where generating novel content, simulating scenarios, or creating synthetic data is essential.

Key Characteristics:

- Generates new content and ideas, including text, images, videos, and music.
- Utilizes deep-learning models trained on vast datasets.

Importance of Generative AI

1. Transforming Customer Experiences:

- Applications like ChatGPT enhance customer interactions, create new types of applications, and boost productivity.

2. Accelerating Research:

- Explores and analyzes complex data in innovative ways.
- Summarizes content, outlines solutions, and creates documentation from research notes.
- Used in the pharma industry for protein sequence generation and drug discovery.

3. Enhancing Customer Experience:

- AI-powered chatbots and virtual assistants provide natural responses and personalized workflows.
- Improves engagement through curated offers and personalized communication.

4. Optimizing Business Processes:

- Applies machine learning across all business lines.
- Extracts and summarizes data, evaluates cost-reduction scenarios, and generates synthetic data for ML processes.

5. Boosting Employee Productivity:

- Augments workflows by generating prototypes, optimizing designs, suggesting code, and producing marketing content.

Impact on Various Domains

- Creative Industries:

- Supports art generation, music composition, and literature creation.
- Healthcare:
 - Facilitates medical imaging analysis, drug discovery, and patient data synthesis.
- Finance:
 - Enhances risk assessment, market prediction, and algorithmic trading.
- Education:
 - Supports personalized learning tools and content creation.

Industry-Specific Benefits

1. Financial Services:

- Improves customer service through chatbots and personalized product recommendations.
- Speeds up loan approvals and detects fraud.
- Provides personalized financial advice at low costs.

2. Healthcare and Life Sciences:

- Accelerates drug discovery with novel protein sequence creation.
- Designs synthetic gene sequences for synthetic biology and biomanufacturing.
- Creates synthetic patient data for training AI models and simulating clinical trials.

3. Automotive and Manufacturing:

- Optimizes mechanical part designs and enhances in-vehicle experiences.
- Improves customer service and manufacturing processes.
- Generates synthetic data for testing applications.

4. Media and Entertainment:

- Produces content efficiently, enhancing music, movies, and games.

- Offers personalized content and advertisements.
- Allows gamers to create avatars and game elements.

5. Telecommunication:

- Reinvents customer experience with conversational agents.
- Optimizes network performance and personalizes customer relationships.

6. Energy:

- Analyzes raw data for usage patterns, product offerings, and efficiency programs.
- Enhances grid management, site safety, and energy production optimization.

Predictive AI and Generative AI

Generative AI and predictive AI are two different types of artificial intelligence (AI) that are used for different purposes. Generative AI is used to create new content, such as music, images, and texts, while predictive AI is used for clustering, classification and regression that often relies on supervised learning and historical training sets that create models for predictions about future states or events. For example, a predictive AI model that is trained on a dataset of historical data about the stock market may be able to make predictions about the future prices of stocks.

Generative AI models are built by training on a large dataset of general examples (such as wikipedia, commoncrawl, etc) and then using that knowledge to generate new examples that are similar to the training data. So the idea is to generate new data. A generative AI model that is trained on a dataset of images of cats may be able to generate new images of cats that are similar to the training data.

Generative AI models are often used in creative applications, such as creating new works of art or music. Predictive AI models are often used in business applications, such as predicting customer behavior or making financial decisions.

Both generative AI and predictive AI models have the potential to improve our lives in many ways. Generative AI models can be used to create new products and services, to automate tasks, and to improve the efficiency of businesses. Predictive AI models can be used to make better decisions, to improve customer service, and to prevent crimes.

However, both generative AI and predictive AI models also have the potential to be used for harmful purposes. Generative AI models can be used to create fake news or to spread misinformation. Predictive AI models can be used to discriminate against people or to invade their privacy.

Difference between generative and predictive AI:

Generative AI:

- Focuses on creating new content, such as music, images, and texts, using deep learning algorithms, often through the emergent behavior of large language models (LLM).
- Involves training on a large dataset of general examples (such as Wikipedia, CommonCrawl, etc.) and using that knowledge to generate new examples that are similar to the training data.
- Used in creative applications, such as creating new works of art or music.
- Requires a reproducible, scalable and responsible approach that minimizes risks of AI safety, misuse mitigation, and model robustness.
- Requires data collection, preparation, training, evaluation, deployment, and monitoring, all encompassed with an MLOps pipeline, but with specific nuances and augmentations needed to adopt generative AI in the enterprise.

Predictive AI:

GenAI

- Focuses on making predictions about future states or events, often relying on supervised learning and historical training sets that create models for clustering, classification, and regression.
- Involves training on a dataset of historical data and making predictions based on that data.
- Used in business applications, such as predicting customer behavior or making financial decisions.
- Requires a reproducible, scalable and responsible approach that minimizes risks of AI safety, misuse mitigation, and model robustness.
- Requires data collection, preparation, training, evaluation, deployment, and monitoring, all encompassed with an MLOps pipeline, but with specific nuances and augmentations needed to adopt predictive AI in the enterprise.

Responsible AI

Regardless of whether we are doing Predictive or Generative AI, we need to be incorporating principles of Responsible AI. It is important to use AI models responsibly and to be aware of the potential risks and how to check and mitigate them. We need to incorporate and develop ethical guidelines that minimize AI Safety issues, discourage abuse/misuse of AI models and monitor and guard model robustness for the development and use of these technologies.

As a subfield of AI, Generative AI focuses on creating new content, such as music, images, and texts, using deep learning algorithms, often through the emergent behavior of Large [Language] Models. (LLM). It is a rapidly growing area, with applications ranging from art and entertainment to scientific discovery and marketing content generation.

Working of Generative AI and Foundation Models

I. Foundation Models

Definition and Characteristics: Foundation models (FMs) are advanced machine learning models trained on extensive and diverse data sources. They are designed to handle a wide range of tasks through their broad training and self-supervised learning.

- **Generalized Training:** FMs are trained on large, unlabeled datasets across various domains, enabling them to learn general patterns and relationships.
- **Versatility and Adaptability:** These models can perform diverse tasks such as natural language understanding, text generation, and image recognition. They can generate coherent text, classify images, and create high-resolution images from textual descriptions.
- **Self-Supervised Learning:** FMs use self-supervised learning techniques to create labels from raw data, which makes them flexible and scalable for various applications.

Examples:

- **BERT (Bidirectional Encoder Representations from Transformers):** Developed by Google, BERT processes text bidirectionally to improve understanding and performance in tasks like question answering and sentence prediction.
- **GPT Series (Generative Pre-trained Transformer):** OpenAI's GPT models, including GPT-3 and GPT-4, are known for their text generation capabilities and handle complex language tasks through extensive parameterization.
- **Stable Diffusion:** A text-to-image model by Stability AI that generates high-definition images from textual prompts using a diffusion process.

Importance: Foundation models have transformed machine learning by enabling more efficient and cost-effective application development. They reduce the need for extensive, task-specific training by providing a robust starting point.

Applications:

1. **Customer Support:** AI-driven chatbots and virtual assistants enhance user experience with accurate and contextually relevant responses.
2. **Language Translation:** Converts text between languages or translates natural language into code.
3. **Content Generation:** Automatically generates articles, blog posts, and other written content based on prompts.
4. **Image Classification and Creation:** Identifies objects within images and generates high-resolution images from text descriptions.
5. **Document Extraction:** Extracts relevant information from documents for analysis or further processing.
6. **Healthcare:** Assists in medical diagnoses, drug discovery, and personalized treatment plans.
7. **Autonomous Vehicles:** Improves navigation systems and decision-making processes for self-driving cars.

Working Mechanism:

- **Neural Networks:** FMs use layers of interconnected nodes to process data. Transformers, for instance, employ attention mechanisms to enhance understanding and generation capabilities.
- **Self-Supervised Learning:** Models generate labels from data itself, allowing for efficient learning from large datasets.
- **Prediction Techniques:** For text, FMs predict the next word based on context. For images, they refine image details using advanced algorithms.

Use Cases:

1. **Language Processing:** Generating coherent text, answering questions, and translating languages.
2. **Visual Comprehension:** Identifying objects, generating images from text, and performing photo and video editing.
3. **Code Generation:** Writing and debugging computer code based on natural language instructions.

4. **Human-Centered Engagement:** Supporting decision-making in areas like clinical diagnoses and analytics, and developing new applications by fine-tuning existing models.
5. **Speech to Text:** Transcribing spoken language into written text and generating captions for videos.

Challenges:

1. **Infrastructure Requirements:** Building and training FMs require substantial computational resources and time.
2. **Front-End Development:** Integrating FMs into practical applications involves additional tools for prompt engineering and fine-tuning.
3. **Comprehension Limitations:** FMs may struggle with understanding complex contexts and lack social or psychological awareness.
4. **Unreliable Answers:** Potential for generating incorrect or inappropriate responses.
5. **Bias and Ethical Considerations:** FMs may inherit biases from training data, necessitating careful data management and mitigation of harmful effects.

II. Large Language Models (LLMs)

Definition and Characteristics: Large Language Models (LLMs) are a subset of foundation models focused on language-based tasks. They are distinguished by their vast number of parameters and capability to handle a broad range of language-related tasks.

- **High Parameter Count:** LLMs like GPT-3 and GPT-4 have billions or even trillions of parameters, enabling them to capture intricate language patterns.
- **Versatility:** They perform multiple language-related tasks, including text generation, summarization, classification, and open-ended conversation.
- **Contextual Understanding:** LLMs leverage extensive pretraining on diverse data to generate and understand content from minimal input.

Examples:

- **GPT-3:** Developed by OpenAI, GPT-3 uses 175 billion parameters to generate coherent and contextually relevant text. It handles tasks from writing articles to answering complex questions.
- **GPT-4:** An evolution of GPT-3, offering enhanced capabilities with improved performance in various language tasks.

Applications:

1. **Text Generation:** Producing human-like text for articles, stories, and creative content based on prompts.
2. **Summarization:** Condensing lengthy documents or articles into concise summaries while retaining key information.
3. **Classification:** Categorizing text into predefined categories, such as sentiment analysis or topic classification.
4. **Open-Ended Conversation:** Engaging in natural language dialogues with users, maintaining context and providing relevant responses.
5. **Information Extraction:** Identifying and extracting specific information from text, such as named entities, dates, or key facts.

III. Overview of Generative Models

Definition and Types: Generative models are algorithms designed to generate new data samples that resemble the data they were trained on. These models are crucial for tasks that involve creating new instances or simulating data based on learned patterns.

Key Types:

1. **Generative Adversarial Networks (GANs):**
 - **Components:** Comprise a generator and a discriminator that compete with each other. The generator creates new data samples, while the discriminator evaluates their authenticity.
 - **Applications:** Used in image and video generation, style transfer (e.g., converting photos to painting styles), and data augmentation (creating synthetic data).

- **Example:** DeepFake technology uses GANs to create realistic synthetic videos of people.

2. Variational Autoencoders (VAEs):

- **Components:** Learn to encode input data into a latent space and then decode it back. They focus on learning meaningful representations of data.
- **Applications:** Useful for generating images, exploring variations of data, and learning complex data distributions.
- **Example:** VAEs generate new faces by learning the underlying distribution of facial features.

3. Transformer Models:

- **Components:** Use attention mechanisms to process and generate sequences of data, excelling in handling long-range dependencies and context.
- **Applications:** Primarily used in natural language processing (NLP) tasks such as text generation, translation, summarization, and question answering.
- **Example:** Models like GPT-3 and BERT use transformer architecture for state-of-the-art results in various NLP tasks.

I. AI vs ML vs Generative AI

Artificial Intelligence (AI):

- **Definition:** AI is a broad field that encompasses a variety of techniques and systems designed to perform tasks that typically require human intelligence. This includes tasks such as problem-solving, reasoning, learning, and understanding natural language.
- **Scope:** AI covers a wide range of applications, from expert systems and robotics to computer vision and natural language processing (NLP).

Machine Learning (ML):


- **Definition:** ML is a subset of AI focused on the development of algorithms that enable systems to learn from and make predictions or decisions based on data.


- **Techniques:** Includes supervised learning (training with labeled data), unsupervised learning (finding hidden patterns in unlabeled data), and reinforcement learning (learning through interaction with an environment).
- **Scope:** ML is specifically concerned with improving the performance of algorithms over time as they are exposed to more data.

Generative AI:

- **Definition:** Generative AI is a specialized subset of AI that focuses on algorithms capable of generating new content or data based on learned patterns from existing data.
- **Techniques:** Uses models such as Generative Adversarial Networks (GANs) and Variational Autoencoders (VAEs) to create new instances of data that resemble training data.
- **Scope:** Primarily used for creating content, such as images, text, and audio, and for tasks that involve generating new data rather than just predicting or classifying existing data.

Generative AI Project Lifecycle





Data Collection and Preparation

- **Data Collection:** Collect data from various sources including internet, social media, sensors, and optimally data warehouses like Google BigQuery or Google Cloud Storage. When collecting data to train generative AI models, ensure that the dataset represents a diverse range of perspectives, backgrounds, and sources. Be cautious about relying solely on data from a single source or domain, as it may introduce bias.
- **Data Creation:** Optionally generate/use synthetic data to augment existing datasets to achieve better performance metrics on trained models.
- **Data Curation:** Clean and organize data through filtering, transformation, integration and labeling of data for supervised learning. Programmatic labeling may be an indispensable part of how we can accelerate training and fine-tuning.
- **Feature Storage:** Store the gold standard of features that have been painstakingly curated and prepped. We do this primarily for AI governance and consistency across the enterprise and the may project that will be attached to it; we will want to store and manage features used in training and inference ; often distinguishing between static and dynamic or train time and run time feature

store capabilities. Looking to provide answers to questions from development teams around: “Where did we get the data from? Who touched the data? Did they bias the data during curation/labeling?” For example, you can use the [Vertex AI Feature Store](#) to support this task.

- **Data Bias Check.** This is an essential step in the machine learning lifecycle, especially for generative AI models, to ensure that the training data and the generated content is fair, unbiased, and representative of the underlying data distribution. Bias in generative AI models can lead to the creation of outputs that reinforce stereotypes, discriminate against certain groups, or misrepresent reality. We can address data bias in predictive and generative AI, right after collecting and curating the data. By incorporating *data bias checks* into the machine learning lifecycle for generative AI, we can work towards creating models that generate fair, unbiased, and representative content, ultimately ensuring that the AI systems we develop align with ethical guidelines and societal values.



Data Inspection: Examine the collected data to identify potential biases or imbalances. This can be done through exploratory data analysis, visualization, or statistical tests. Analyze the data for potential issues, such as underrepresentation of certain groups, overrepresentation of specific topics, or skewed sentiment distributions.

Bias Mitigation Techniques: Apply various techniques to mitigate the identified biases in the data. These techniques may include re-sampling, re-weighting, or generating synthetic examples to balance the dataset. Additionally, consider using de-biasing algorithms during the model training process to minimize the impact of biases on the model’s output.

Model Evaluation: Develop evaluation metrics and techniques that specifically assess the model’s fairness and bias. These might include measuring the model’s performance across different demographic groups, evaluating the diversity of

generated samples, or analyzing the model's behavior when given prompts that are prone to biased outputs.

Continuous Monitoring: After deploying the generative AI model, monitor its performance and outputs to ensure that biases do not emerge over time. Set up feedback loops to gather user feedback and input, which can be used to identify and address any newly discovered biases.

Transparency and Accountability: Document and share the steps taken to address data bias in the generative AI model. Communicate the model's limitations and potential biases to end-users and stakeholders, fostering trust and setting realistic expectations for the model's performance.

Model Training & Experimentation

- **Experimentation:** This is name of the game in machine learning.
- **Pre-trained model training:** Train a model using a pre-trained model as a starting point
- **Fine-tune models:** Update an existing model with new data or for a new task
- **Model registry:** This is a means and mechanism for Model Governance and Model Risk Management. To use the Model Registry we would Store and manage models for versioning, reuse, and auditing, marking the meta-data of the model to see where the data came from (Data Provenance, Lineage), how the model was trained, etc.

Distributed Training

- **Data parallelism:** Split data across multiple devices and training them simultaneously to speed up training
- **Model parallelism:** Split the model across multiple devices and training them simultaneously to fit larger models into memory

- Federated learning: Train models collaboratively across multiple devices or organizations while preserving privacy and security.

Model Adaptation

Transfer learning

Adapt a pre-trained model to a new domain or task

Fine-tune models

Update an existing model with new data or for a new task

AI Safety Mitigation

Ensure the model is safe, ethical, and trustworthy by mitigating misuse, improving robustness, and ensuring privacy and security.

Distillation

Train or create a much smaller model with similar performance and downstream task capabilities by using a student/teacher few-shot learning approach to train a much smaller model, or, prune a model to arrive at a smaller footprint deployable model.

There are several methods to distill large language models (LLMs) to arrive at significantly smaller models with comparable performance. Let's explore some common approaches .

1. **Pruning:** This involves removing unimportant weights or neurons from the LLM while preserving its accuracy. Pruning can be done using various techniques such as magnitude-based pruning, sensitivity-based pruning, and weight-rewarding.
2. **Quantization:** This involves reducing the precision of the LLM's weights and activations. For example, weights that were initially represented as 32-bit

floating-point numbers can be quantized to 8-bit integers. This reduces the model's memory footprint without significantly impacting its performance.

3. **Knowledge Distillation:** This involves training a smaller model to mimic the output of the LLM on a set of training examples. The idea is to transfer the knowledge from the larger model to the smaller one. This approach can be used to train models that are orders of magnitude smaller than the original model with only a small reduction in accuracy.
4. **Low-rank Factorization:** This involves decomposing the weight matrices of the LLM into low-rank matrices. This reduces the number of parameters in the model while preserving its accuracy.
5. **Compact Embeddings:** This involves reducing the dimensionality of the input and output embeddings of the LLM. This reduces the number of parameters in the model and speeds up inference.
6. **Architectural Changes:** This involves changing the architecture of the LLM to make it more efficient. For example, the transformer architecture can be modified to reduce its memory footprint or the number of attention heads can be reduced.
7. **Parameter Sharing:** This involves sharing parameters between different parts of the model. For example, the same weight matrix can be used for multiple layers in a neural network, reducing the number of parameters and improving the model's efficiency.



Depending on the specific use case, a combination of these techniques may be used to achieve the desired level of performance and efficiency.


Another approach is the Teacher-student few-shot approach to producing smaller models. This approach is actually a variant of knowledge distillation, but it specifically focuses on few-shot learning — the ability to learn from a small number of examples.

In the teacher-student [1] few-shot approach, a large LLM (the teacher model) is used to generate synthetic examples, which are used to train a smaller LLM (the student

model). The idea is to use the teacher model to generate examples that are representative of the underlying distribution of natural language, even though the student model is only trained on a small number of examples.

The teacher-student process involves the following steps:

1. Train a large LLM (the teacher model) on a large corpus of text.
2. Generate synthetic examples using the teacher model by either sampling from the model's distribution or by using specific prompts.
3. Train a smaller LLM (the student model) on a small number of labeled examples, along with the synthetic examples generated by the teacher model.
4. Fine-tune the student model on a small validation set to optimize its performance.



The teacher-student few-shot approach can significantly reduce the number of examples required to train a high-quality LLM. This is particularly useful in scenarios where labeled examples are scarce or when fine-tuning is required for specific domains or tasks. However, it should be noted that the quality of the synthetic examples generated by the teacher model can impact the performance of the student model. Therefore, it is important to carefully evaluate the quality of the synthetic examples before using them to train the student model.

Model Deployment

Let's discuss how to serve the model using [Google Cloud's Vertex AI](#) endpoint. Serving a model on a Vertex AI endpoint involves packaging the model into a container image, uploading the image to a container registry, creating a model resource and endpoint in the Vertex AI console, deploying the model to the endpoint, testing the endpoint, and monitoring its performance.

1. **Prepare the Model.** First, you need to prepare the trained model for deployment. This typically involves packaging the model and any required dependencies into a container image, such as a Docker image.
 2. **Upload the Container Image.** Next, you need to upload the container image to a container registry, such as Google Container Registry. This will allow you to easily deploy the image to a Vertex AI endpoint.
 3. **Create a Model Resource.** In the Vertex AI console, create a model resource that will represent your deployed model. This involves specifying the model's name, description, and any other metadata that may be required.
 4. **Create an Endpoint.** Create a Vertex AI endpoint to serve your model. This involves specifying the endpoint's name, description, and the region where it will be deployed.
 5. **Deploy the Model.** Once the endpoint is created, you can deploy your model to the endpoint. This involves specifying the container image to use and any required settings, such as the number of replicas to deploy.
- Edge computing. Deploy models on edge devices for low-latency, privacy-preserving, and real-time inference. Distill models for edge compute as LLMs or Foundation Models in general are huge and will not fit in the footprint of small devices.
 - Cloud deployment. Deploy models on cloud platforms for scalability, accessibility, and cost-effectiveness



5.1 Test the Endpoint. Once the model is deployed, you can test the endpoint by sending sample requests to it. This will allow you to verify that the endpoint is serving the model correctly and that the model is generating accurate predictions.

5.2 Monitor the Endpoint. After the endpoint is deployed, you should monitor its performance to ensure that it is meeting the required service level agreements (SLAs). This may involve monitoring latency, throughput, and error rates, and making

adjustments as needed. Monitoring and governance. Monitor model performance, detecting issues, and enforcing compliance with regulations and policies.

6. Serving

Here are some considerations for serving the model or running inference/ generating content :

1. **Prepare the Input Data.** Once the model is loaded, the next step is to prepare the input data that will be fed into the model. This may involve preprocessing the input data, such as converting text to numerical embeddings or resizing images.
2. **Feed Input to the Model.** After the input data is prepared, it can be fed into the loaded model for inference. This typically involves calling the model's `predict` or `generate` method and passing in the input data.
3. **Generate Output.** The model will then generate output based on the input data. The output can take many forms depending on the specific task and model architecture. For example, a generative text model may generate a sequence of words or sentences, while a generative image model may generate an image or a sequence of images.
4. **Postprocess the Output.** After the output is generated, it may require postprocessing to make it usable for downstream tasks. This may involve converting the output from numerical embeddings to text, resizing images, or applying other transformations.

Model Maintenance

- Data and model drift. Monitor and managing changes in data distribution and model performance over time
- Model versioning. Manage different versions of the model for traceability, reproducibility, and comparison

- Prompt-tuning, Prompt design: Optimize prompts used in language generation models to improve quality, coherence, and relevance.

Performance optimization

- Hardware acceleration. Optimize the model for specific hardware accelerators to improve inference speed and efficiency
- Quantization. Reduce the model size and computation requirements by converting it to lower precision data types
- Pruning. Reduce the model size and computation requirements by removing unimportant weights and neurons.



Prompt-tuning, Prompt Design, Instruction-tuning

It is important to compare and contrast the various ways in which we can tune or customize the models.

Prompt design. In this phase, you design the initial prompt that will be used to generate responses from your language model. The prompt should be carefully crafted to elicit the desired type of response and to provide enough context for the model to generate coherent and relevant output.

Prompt tuning. Once you have a draft prompt, you'll need to fine-tune it to improve the quality of the responses generated by your language model. This might involve tweaking the wording of the prompt, adjusting the level of detail provided, or adding additional context or constraints to guide the model's output.

Instruction tuning. In addition to tuning the prompt itself, you may also need to provide additional instructions or guidance to the model to help it generate the type of output you're looking for. This might involve specifying certain criteria for the output, providing sample outputs to use as references, or setting constraints on the content or style of the generated text.


Fine-tuning. Once you have a well-designed prompt and clear instructions for your model, you'll need to fine-tune the model itself to optimize its performance on your specific task. This involves training the model on a dataset of examples and adjusting its parameters to minimize the loss function and improve the quality of the generated text. The goal of fine-tuning is to create a model that can consistently generate high-quality responses that meet the criteria specified in the prompt and instructions.

Variation 1: Prompt-first GenAI

Drastically Reduce Application development time with Prompt Based Human Interaction

One of the ways in which generative AI is productive is the way in which we can employ it in the early stages of the software development lifecycle. Through rapidly prototyping synthetic data sets to code generation snippets we can accelerate the process by using prompt design to produce faster turn around and break through initial barriers of doubt to arrive at definitive solution paths : I will use this type of data with this type of algorithm to accomplish this specific task.

Thus, you can almost think of the Generative AI Lifecycle as starting with prompt design or tuning and gather data in parallel or after the initial prototype has been determined to be feasible.



Prompt design and tuning can be incorporated into the early stages of the GenAI lifecycle to optimize the AI model's performance and effectiveness even before the data preparation stage. This can help ensure that the AI model aligns better with the project goals and also provides a set of more accurate and relevant responses to subsequent inputs, prompts and inferences. The relevant tasks within this early stage consists of the following.

The GenAI lifecycle prompt-based variation gives us a proactive approach to prompt optimization that can lead to much more efficient AI development lifecycle and resulting models. In this way, we can better align with the desired goals and provide improved performance on specific lifecycle tasks or industry domains.

Problem formulation and planning

Before starting with data preparation, it's essential to define the problem you want the AI model to solve. During this stage, you can identify the specific tasks or instructions the model should be capable of handling. This helps set the stage for designing appropriate prompts and planning for potential tuning strategies later on.

Preliminary research and exploration

During this stage, you can explore existing AI models and solutions to understand their strengths, weaknesses, and limitations. This can give you insights into which prompt design strategies or tuning techniques might be effective for your specific problem. Also, you can gather information on potential pitfalls and challenges that may need to be addressed during prompt design or tuning.

Designing and prototyping prompts

Instead of waiting until the data preparation stage, you can start prototyping potential prompts early in the process. This enables you to experiment with different phrasing, context, and formatting to test the AI model's responsiveness and effectiveness. It's essential to iterate and refine your prompt design to optimize the AI model's performance at this stage.

Feedback and collaboration

Collaborating with domain experts or target users during the early stages of the ML lifecycle can help you gain insights into the specific needs and requirements of your AI model. This feedback can be invaluable in shaping your prompt design or tuning strategy and ensuring the AI model's relevance and effectiveness in its intended domain.

The Pros and Cons of Generative AI

Advantages:

1. Increased Productivity:

- **Automation:** Generative AI can automate repetitive tasks, leading to increased efficiency and faster completion of processes.
- **Speed:** By accelerating tasks such as content generation and data analysis, it can significantly reduce the time required for various applications.

2. Lower Skill and Time Barriers:

- **Accessibility:** It democratizes content creation and creative applications by lowering the barriers for individuals who may not have specialized skills or extensive time resources.

3. **Complex Data Analysis:**

- **Exploration:** Enables the analysis and exploration of complex datasets, providing insights and facilitating tasks that may be challenging to handle manually.

4. **Synthetic Data Creation:**

- **Training and Improvement:** Can generate synthetic data that is used to train and enhance other AI systems, improving their performance and robustness.

Disadvantages:

1. **Hallucination:**

- **Generation of Nonsense:** Generative AI models may produce outputs that are nonsensical, factually incorrect, or lack common-sense logic, leading to misinformation or erroneous results.

2. **Reliance on Data Labeling:**

- **Quality and Veracity Issues:** Even though many models can be trained in an unsupervised manner, the quality of data and its labeling remain critical. The reliance on low-paid workers for data labeling and enrichment can impact the quality of training data.

3. **Difficulty with Content Moderation:**

- **Inappropriate Content:** AI models may struggle to effectively recognize and filter out inappropriate or harmful content. Much of this task still requires human oversight, which can be resource-intensive.

4. **Ethical Issues:**

- **Bias and Discrimination:** Generative AI can replicate or amplify existing biases in training data, leading to discriminatory outcomes. For instance, biased AI recruiting tools have shown negative impacts on certain groups.

5. **Legal and Regulatory Issues:**

- **Copyright Concerns:** The use of large datasets for training AI models can lead to copyright issues regarding the materials used and the generated outputs.
- **Privacy:** Generative AI raises questions about data privacy, including how personal and business-related data is collected, stored, and used.
- **Autonomy and Responsibility:** Determining liability in cases involving autonomous systems (e.g., self-driving cars) poses challenges, especially when accidents or errors occur.

6. Political Implications:

- **False Information:** Generative AI can create misleading or false information, impacting media and public perception. It may also be used to interfere with democratic processes by generating large volumes of deceptive content.

7. Energy Consumption:

- **Ecological Impact:** The computational power required for training and running AI models results in significant energy consumption, contributing to environmental concerns.

Explainable AI

Explainable artificial intelligence (XAI) is a set of processes and methods that allows human users to comprehend and trust the results and output created by machine learning algorithms.

Explainable AI is used to describe an AI model, its expected impact and potential biases. It helps characterize model accuracy, fairness, transparency and outcomes in AI-powered decision making. Explainable AI is crucial for an organization in building trust and confidence when putting AI models into production. AI explainability also helps an organization adopt a responsible approach to AI development.

As AI becomes more advanced, humans are challenged to comprehend and retrace how the algorithm came to a result. The whole calculation process is turned into what is commonly referred to as a “black box” that is impossible to interpret. These

black box models are created directly from the data. And, not even the engineers or data scientists who create the algorithm can understand or explain what exactly is happening inside them or how the AI algorithm arrived at a specific result.

Importance of explainable AI

It is crucial for an organization to have a full understanding of the AI decision-making processes with model monitoring and accountability of AI and not to trust them blindly. Explainable AI can help humans understand and explain machine learning (ML) algorithms, deep learning and neural networks.

ML models are often thought of as black boxes that are impossible to interpret.² Neural networks used in deep learning are some of the hardest for a human to understand. Bias, often based on race, gender, age or location, has been a long-standing risk in training AI models. Further, AI model performance can drift or degrade because production data differs from training data. This makes it crucial for a business to continuously monitor and manage models to promote AI explainability while measuring the business impact of using such algorithms. Explainable AI also helps promote end user trust, model auditability and productive use of AI. It also mitigates compliance, legal, security and reputational risks of production AI.

Explainable AI is one of the key requirements for implementing responsible AI, a methodology for the large-scale implementation of AI methods in real organizations with fairness, model explainability and accountability.³ To help adopt AI responsibly, organizations need to embed ethical principles into AI applications and processes by building AI systems based on trust and transparency.

Working of explainable AI

With explainable AI – as well as interpretable machine learning – organizations can gain access to AI technology's underlying decision-making and are empowered to make adjustments. Explainable AI can improve the user experience of a product or service by helping the end user trust that the AI is making good decisions. When

do AI systems give enough confidence in the decision that you can trust it, and how can the AI system correct errors that arise?⁴

As AI becomes more advanced, ML processes still need to be understood and controlled to ensure AI model results are accurate. Let's look at the difference between AI and XAI, the methods and techniques used to turn AI to XAI, and the difference between interpreting and explaining AI processes.

Comparing AI and XAI

XAI implements specific techniques and methods to ensure that each decision made during the ML process can be traced and explained. AI, on the other hand, often arrives at a result using an ML algorithm, but the architects of the AI systems do not fully understand how the algorithm reached that result. This makes it hard to check for accuracy and leads to loss of control, accountability and auditability.

Explainable AI techniques

The setup of XAI techniques consists of three main methods. Prediction accuracy and traceability address technology requirements while decision understanding addresses human needs. Explainable AI — especially explainable machine learning — will be essential if future warfighters are to understand, appropriately trust, and effectively manage an emerging generation of artificially intelligent machine partners

Prediction accuracy

Accuracy is a key component of how successful the use of AI is in everyday operation. By running simulations and comparing XAI output to the results in the training data set, the prediction accuracy can be determined. The most popular technique used for this is Local Interpretable Model-Agnostic Explanations (LIME), which explains the prediction of classifiers by the ML algorithm.

Traceability

Traceability is another key technique for accomplishing XAI. This is achieved, for example, by limiting the way decisions can be made and setting up a narrower

scope for ML rules and features. An example of a traceability XAI technique is DeepLIFT (Deep Learning Important FeaTures), which compares the activation of each neuron to its reference neuron and shows a traceable link between each activated neuron and even shows dependencies between them.

Decision

understanding

This is the human factor. Many people have a distrust in AI, yet to work with it efficiently, they need to learn to trust it. This is accomplished by educating the team working with the AI so they can understand how and why the AI makes decisions.

Explainability versus interpretability in AI

Interpretability is the degree to which an observer can understand the cause of a decision. It is the success rate that humans can predict for the result of an AI output, while explainability goes a step further and looks at how the AI arrived at the result.

How does explainable AI relate to responsible AI?

Explainable AI and responsible AI have similar objectives, yet different approaches. Here are the main differences between explainable and responsible AI:

- Explainable AI looks at AI results after the results are computed.
- Responsible AI looks at AI during the planning stages to make the AI algorithm responsible before the results are computed.
- Explainable and responsible AI can work together to make better AI.

Continuous model evaluation

With explainable AI, a business can troubleshoot and improve model performance while helping stakeholders understand the behaviors of AI models. Investigating model behaviors through tracking model insights on deployment status, fairness, quality and drift is essential to scaling AI.

Continuous model evaluation empowers a business to compare model predictions, quantify model risk and optimize model performance. Displaying positive and negative values in model behaviors with data used to generate explanation speeds

model evaluations. A data and AI platform can generate feature attributions for model predictions and empower teams to visually investigate model behavior with interactive charts and exportable documents.

Benefits of explainable AI

Operationalize AI with trust and confidence

Build trust in production AI. Rapidly bring your AI models to production. Ensure interpretability and explainability of AI models. Simplify the process of model evaluation while increasing model transparency and traceability.

Speed time to AI results

Systematically monitor and manage models to optimize business outcomes. Continually evaluate and improve model performance. Fine-tune model development efforts based on continuous evaluation.

Mitigate risk and cost of model governance

Keep your AI models explainable and transparent. Manage regulatory, compliance, risk and other requirements. Minimize overhead of manual inspection and costly errors. Mitigate risk of unintended bias.

Five considerations for explainable AI

To drive desirable outcomes with explainable AI, consider the following.

Fairness and debiasing: Manage and monitor fairness. Scan your deployment for potential biases.

Model drift mitigation: Analyze your model and make recommendations based on the most logical outcome. Alert when models deviate from the intended outcomes.

Model risk management: Quantify and mitigate model risk. Get alerted when a model performs inadequately. Understand what happened when deviations persist.

Lifecycle automation: Build, run and manage models as part of integrated data and AI services. Unify the tools and processes on a platform to monitor models and share outcomes. Explain the dependencies of machine learning models.

Multicloud-ready: Deploy AI projects across hybrid clouds including public clouds, private clouds and on premises. Promote trust and confidence with explainable AI.

Use cases for explainable AI

- **Healthcare:** Accelerate diagnostics, image analysis, resource optimization and medical diagnosis. Improve transparency and traceability in decision-making for patient care. Streamline the pharmaceutical approval process with explainable AI.
- **Financial services:** Improve customer experiences with a transparent loan and credit approval process. Speed credit risk, wealth management and financial crime risk assessments. Accelerate resolution of potential complaints and issues. Increase confidence in pricing, product recommendations and investment services.
- **Criminal justice:** Optimize processes for prediction and risk assessment. Accelerate resolutions using explainable AI on DNA analysis, prison population analysis and crime forecasting. Detect potential biases in training data and algorithms.

The below diagram details the lifecycle in terms of time and effort required.

	Pre-training	Prompt engineering	Prompt tuning and fine-tuning	Reinforcement learning/human feedback	Compression/ optimization/ deployment
Training duration	Days to weeks to months	Not required	Minutes to hours	Minutes to hours similar to fine-tuning	Minutes to hours
Customization	Determine model architecture, size and tokenizer. Choose vocabulary size and # of tokens for input/context Large amount of domain training data	No model weights Only prompt customization	Tune for specific tasks Add domain-specific data Update LLM model or adapter weights	Need separate reward model to align with human goals (helpful, honest, harmless) Update LLM model or adapter weights	Reduce model size through model pruning, weight quantization, distillation Smaller size, faster inference
Objective	Next-token prediction	Increase task performance	Increase task performance	Increase alignment with human preferences	Increase inference performance
Expertise	High	Low	Medium	Medium-High	Medium

key additions to the traditional AI lifecycle when adopting generative AI:

1. [Prompt-based Prototyping]. In this variation of the starting phase, before data preparation in the GenAI lifecycle, we proactively focus on designing and testing candidate prompts for an AI language model. This method allows experimentation with data sets, formats, various input phrasings, contexts, and formats to evaluate the model's performance in addressing desired tasks or goals. Iterating and refining prompts based on performance helps identify and resolve potential issues in the model's comprehension and output. By implementing prompt-based prototyping early, you can better align the AI model with its intended purpose, ultimately leading to more effective data preparation strategies.
2. Data Generation and Programmatic Data Labelling: In addition to data collection and preparation, generative AI requires the usage or creation of large datasets for training generative models. This involves generating synthetic data or using unsupervised learning techniques to learn patterns from existing, unlabelled data.

3. Prompt- , instruction- and fine-tuning. This has huge implications for model usage and human computer interaction for higher efficiency and productivity, using GenAI as a human augmentation mechanism.
4. Model Selection: With so many models in multiple Model Gardens and hubs, which open to choose?
5. Model Evaluation: Traditional model evaluation techniques such as accuracy, precision, and recall may not be suitable for evaluating generative models. New evaluation metrics need to be developed, such as sample diversity, quality, and consistency.
6. Interpretability: Generative models can be difficult to interpret, and understanding how they generate outputs is important for ensuring their responsible use. Techniques such as visualization, attribution, and counterfactual analysis can be used to improve interpretability.
7. Responsible AI & Ethical Considerations: Generative AI can raise ethical concerns related to bias, fairness, privacy, and safety. These concerns need to be addressed through the development of ethical guidelines, transparency, and accountability mechanisms.
8. Distillation: Creating smaller deployable models for cost and edge computing restrictions.
9. Deployment: Generative AI models may have different deployment requirements than traditional predictive models. They may require specialized hardware, longer inference times, and different deployment architectures.
10. Tuning: even after deployment, tuning the deployed model to structure the desired types of responses.
11. Monitoring: Generative models can generate unexpected outputs over time, and monitoring their performance is critical for ensuring their ongoing reliability and safety. Techniques such as drift detection and feedback loops can be used to monitor and improve model performance.

References

[1] Modeling Teacher-Student Techniques in Deep Neural Networks for Knowledge Distillation, [Sajjad Abbasi](#), [Mohsen Hajabdollahi](#), [Nader Karimi](#), [Shadrokh Samavi](#). <https://arxiv.org/abs/1912.13179>

[2] The Pupil Has Become the Master: Teacher-Student Model-Based Word Embedding Distillation with Ensemble Learning,

[Bonggun Shin](#), [Hao Yang](#), [Jinho D. Choi](#). <https://www.ijcai.org/proceedings/2019/0477.pdf>

[3] Language models are few-shot learners, Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., ... & Amodei, D. (2020). arXiv preprint arXiv:2005.14165.

Demonstration

Assignment for explainable AI: <https://pair-code.github.io/what-if-tool/explore/>
<https://aif360.res.ibm.com/>
<https://github.com/marcotcr/lime>
https://github.com/shap/shap/blob/master/notebooks/api_examples/migrating-to-new-api.ipynb
<https://aif360.res.ibm.com/data#>

RSCOE CSBS

RSCOE CSBS